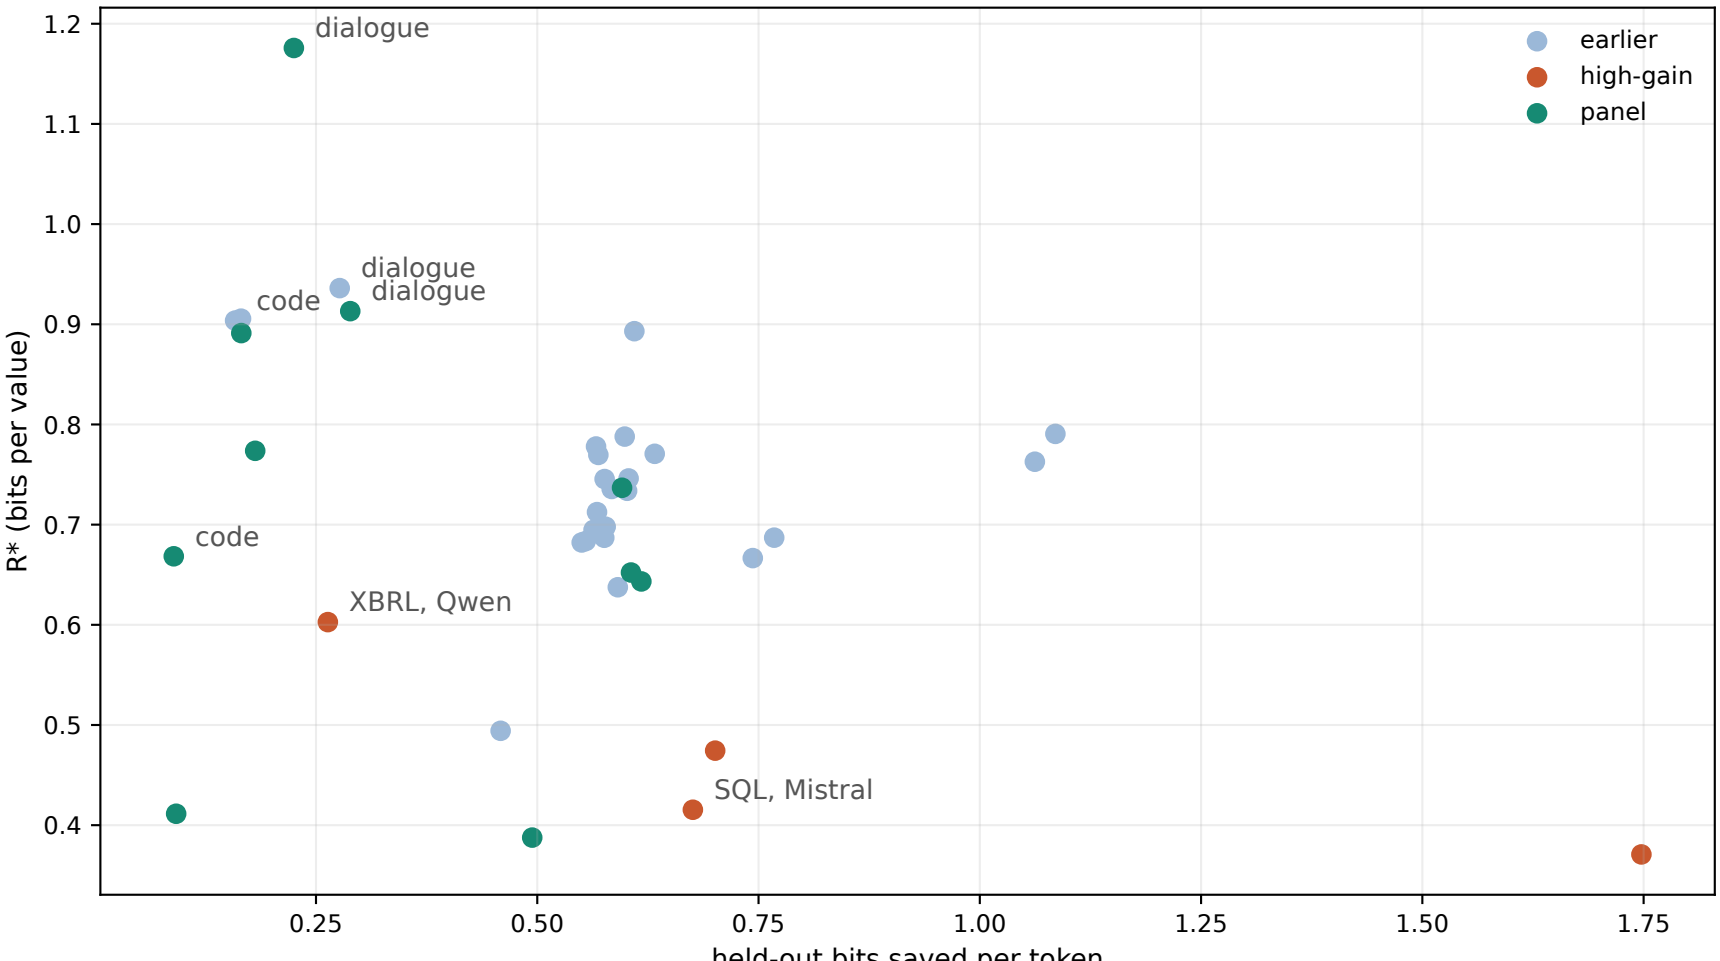

The corpus that changes the model most needs the smallest adapter
38 arms, five corpora, two base models: r = -0.37
earlier
high-gain
panel
dialogue
dialogue
dialogue
code
code
XBRL, Qwen
SQL, Mistral
R* (bits per value)
held-out bits saved per token
1.2
1.1
1.0
0.9
0.8
0.7
0.6
0.5
0.4
0.25
0.50
0.75
1.00
1.25
1.50
1.75