

If we stop after the model predicts the endpoint, delete the original history, and pass only that predicted state to a fresh model call, final accuracy returns to 76.98%. This matches endpoint-inference accuracy and therefore isolates the one-pass deficit to routing or composition rather than state inference. Passing the correct endpoint reaches 100%, as does explicit step-by-step decimal execution, ruling out the final lookup itself as the bottleneck.

Hidden-state probes also retain information about the path used to reach the endpoint, and the late-layer subspace we tested does not outperform a random subspace control. We therefore do not claim that the frozen model exposes one clean native code that can simply be patched between computations.